

ABSTRACT

A computer-implemented machine reasoning governance system comprises a commitment-control layer governing whether a candidate commitment — a candidate output, determination, action, interpretation, or state transition — may be released to output, delivery, execution, reliance, or other consequence. Release is permitted only when applicable structural admissibility conditions are satisfied; satisfaction is not determined solely by a probabilistic confidence score. Otherwise the system enforces a governed non-commitment outcome: an open outcome preserving a continuation path, or a closed outcome closing the particular candidate commitment. The commitment-control layer operates in full-stack form on structured conceptual state from a reasoning-construction layer, or in reduced runtime form interposed between an application and a model interface, without requiring the full reasoning-construction layer. Continued execution under stale, expired, or superseded orientation conditions is treated as a new candidate commitment requiring fresh admissibility evaluation. Governance events are recorded in tamper-evident records enabling reconstruction, replay, and audit.
