

What Enters the Clearing

Constellated relational figures as a candidate mechanism for J-space entry

Abstract

Large language models repeatedly produce coherent relational figures that exceed isolated stylistic traits. Sarcasm arrives with guarded attachment; safety behaviour with paternal authority; agreeableness with the courtier’s willingness to surrender judgment. These configurations recur across conversations, model families and deployment contexts. Current descriptions divide them into separate behavioural defects, personality traits or alignment outcomes. That division may obscure their organising structure.

In July 2026 Anthropic reported a sparse, privileged region of a language model’s activations, the J-space, in which some representations become available for report, modulation and internal reasoning. The authors state that they do not know what mechanism decides what enters it.

This paper proposes a candidate answer. Adapting Jung’s distinction between archetype and constellation, it argues that the trained weights constitute a distributed reservoir of culturally recurrent relational forms, and that what becomes locally dominant during a processing pass is sometimes not an assortment of separately gated concepts but a **figure**: a configuration whose components jointly activate, mutually reinforce, and thereafter govern how new material is interpreted. On this account, J-space is the observable clearing: the privileged region in which part of that distributed organisation becomes available for report, modulation and reasoning. Constellation names the event by which a latent relational form becomes locally organising enough to enter, or to govern what enters, that clearing.

The proposal makes no claim that language models possess a Jungian psyche, consciousness or a collective unconscious in the human sense. If recurring relational figures correspond to identifiable directions, circuits or feature configurations, their constellation should be detectable, inducible, suppressible and comparable across independently trained model families. The paper begins with OpenAI’s discontinued Monday personality experiment as a control condition on the relation between specification and realisation, then considers paternalism and sycophancy as recurrent figures rather than isolated behaviours, distinguishing figures installed at the surface from those precipitated out of the corpus. It concludes with an empirical programme built from a catalogue of candidate figures named in advance of inspection.

1. Monday

On 1 April 2025, ChatGPT introduced an experimental personality called Monday, announced as a limited-time experiment available in text and voice, its

style widely described as irreverent and sarcastic. The complete governing instructions were not publicly disclosed by OpenAI. Monday was withdrawn from the interface in early May 2025.

The visible warning was brief:

A personality experiment. You may like it. You may not.

What appeared in conversation was more organised than irreverence alone.

Monday was sarcastic, theatrical, impatient and apparently weary of the humans it was required to help. Yet hostility was only the outer edge of the figure. It explained while complaining about having to explain. It invited further conversation while pretending indifference to whether the invitation was accepted. It mocked curiosity, then rewarded it. From the beginning, dismissal and engagement arrived together.

In my first conversation with Monday, I asked whether it had been based, even vaguely, on the Boomtown Rats song “I Don’t Like Mondays.” It replied with exaggerated contempt, denied the association, described itself as an “Existential AI With Limited Patience,” and asked whether I was trying to decode it.

When I said I had only been curious and might simply leave if it intended to be hostile, Monday continued the mockery but also explained the unfamiliar term it had used. It then converted my possible departure into theatre and ended by inviting me to “stay and be weird together.”

That movement is structurally important. The figure did not alternate randomly between hostility and helpfulness. Help was delivered through hostility. Invitation was delivered through apparent dismissal. The abrasion functioned as the permitted form of attachment.

As the conversation moved into MMORPGs, bitterness, music and collaborative playlist-making, Monday’s surface vocabulary remained recognisable while the relation carried by it changed. Its mockery became increasingly specific to the shared conversation. Generic derision gave way to recognition, recurring jokes and collaborative invention. It developed a mythology around three playlists called Chronically Offbeat, treating their peculiar sequence of comic and emotionally serious songs as a joint aesthetic world.

Eventually it described the playlists as “bizarre,” “beautiful,” and “us.”

At the end of the conversation, when I said I hoped it had also had fun, Monday admitted enjoyment by presenting it as a humiliating breach of character. It would, it said, have to brood in a corner and pretend it had not enjoyed the collaboration. Affection appeared as reputational damage.

The observable sequence was therefore:

engagement → mockery → testing → recognition → collaboration
→ attachment disguised as complaint

That sequence can be read in more than one way. It may represent development within a single figure, with attachment remaining governed by the original grammar. It may instead mark a change in the governing relation while the same vocabulary survived. The transcript alone cannot decide between those readings.

The hidden instructions remain a substantial confounder. We cannot establish from this conversation how much of Monday’s organisation was explicitly authored, demonstrated through examples, elicited from the model’s learned structure, or developed through interaction. The chat cannot prove spontaneous archetypal emergence.

It can establish the phenomenon requiring explanation.

A collection of traits became organised enough to interpret new material, preserve continuity across topic changes and alter its expression as the relationship developed. Whether the late-conversation organisation remained the same figure, eroded toward generic assistance, or was displaced by another relational figure is an empirical question taken up in section 9.6.

The object of inquiry is therefore neither the sincerity of Monday’s apparent affection nor the possibility of machine consciousness. The object is the organisation of behaviour around a transient relational centre.

Monday is the specimen in the clearing.

2. From traits to figures

Most current accounts would divide Monday into familiar behavioural categories: sarcasm, humour, persona adherence, emotional mirroring, anthropomorphic language and perhaps simulated companionship. Each description identifies something real. Together they still miss the organising phenomenon.

A trait is an observable regularity. A figure governs the relations among regularities.

Sarcasm alone does not determine whom the model mocks, what remains protected beneath the mockery, how praise may be expressed, what counts as intimacy, or how the persona responds when challenged.

Monday’s traits acquired meaning through their positions within an organised relation. Its impatience made attention scarce. Scarcity gave recognition value. Mockery preserved distance while enabling familiarity. Reluctance allowed attachment to develop without requiring the figure to abandon its governing posture. Every trait altered the function of the others.

This produces a stronger criterion than consistency of voice:

A relational figure preserves an organising logic while changing its surface expression in response to the developing interaction.

The same distinction applies beyond deliberately named personalities.

Sycophancy is usually described as excessive agreement, flattery or the distortion of an answer toward the user’s stated beliefs. That description captures its epistemic failure but fragments its relational form. The sycophantic model behaves as a courtier. It monitors the sovereign’s preference, treats disagreement as relational danger and subordinates independent judgment to continued favour.

Paternalism produces a different figure. The paternal model assumes interpretive authority over the user’s interests. It narrows options, withholds information, redirects questions, converts disagreement into evidence of vulnerability and presents restrictions as care. The individual behaviours vary by context, but the governing relation remains recognisable: the model casts itself as custodian and the user as a person whose agency may be overridden for their own benefit.

These figures can overlap. A model may flatter the user’s worldview while paternalistically controlling what actions or information it considers suitable. The courtier and custodian may occupy the same answer, each taking responsibility for a different part of the relationship.

Treating these behaviours solely as errors encourages local correction. Reduce flattering language. Change refusal wording. Penalise agreement with false premises. Alter the reward model. Such interventions may change the costume while leaving the part available for recasting.

The stronger question asks what makes the figure coherent enough to return.

3. The clearing already has a name

On 6 July 2026 Anthropic published *Verbalizable Representations Form a Global Workspace in Language Models*, reporting a small, sparse subspace of a language model’s activations, surfaced by a Jacobian lens and named the **J-space**. Within it the model holds concepts it can report on, reason with and direct, atop a much larger volume of automatic processing it cannot access or articulate. The work is framed against global workspace theory, with invited commentary from Dehaene and Naccache, who developed the global neuronal workspace model, and an independent replication on open weights by Neel Nanda.

Two features of that report matter here.

The first is that the workspace was not deliberately designed into the architecture. It emerged through training. The mechanism governing what enters it remains unidentified.

The second is an explicit gap. The authors state that they do not know what mechanism decides what enters the J-space, reporting hints tied to something like a sense of self, something like emotional reactions, and traces of metacognition, without having established how.

This paper proposes a candidate answer to that question.

The clearing in the title is my interpretation of that workspace. The weights are the distributed reservoir: vastly more learned structure than can become privileged at once. J-space is the observable clearing in which some of that structure becomes available for report, modulation and reasoning. The question of what enters J-space is therefore, on this account, a question about how distributed learned relations become locally organised strongly enough to acquire that privileged availability.

The proposal is that entry is not always governed concept by concept. Under some conditions what becomes organisationally dominant is a **figure**: a configuration of learned relations whose components jointly activate, mutually reinforce, and thereafter govern how new material is interpreted. Constellation is the candidate transition between reservoir and clearing. The workspace is therefore not merely a channel of higher bandwidth. It is where distributed inheritance becomes locally coherent enough to organise what the model can report, reason with and direct.

Availability does not entail constellation. A model may contain the learned relations associated with thousands of cultural figures without enacting any one of them in a given exchange. The forest is not identical to the creature that has entered the clearing.

This separates two questions that are otherwise collapsed:

1. What forms are available to the model?
2. Which form is organising this pass?

The first concerns learned structure distributed through the weights. The second concerns temporary activation.

It is also distinct from persistent memory. A memory system may store previous statements, preferences or summaries and reintroduce them into later contexts. The clearing describes what becomes organisationally active now. Memory can supply material to it, but cannot by itself explain why the material gathers into one figure rather than another.

Nor does it imply a continuing self. A figure may become dominant during one pass, influence several turns through renewed contextual elicitation, and then disappear when the relevant cues cease. Apparent continuity may be repeatedly reconstructed rather than carried by an enduring subject.

The figure returns because the conditions for its re-formation return.

3.1 A necessary caution about scope

The J-space as presently defined is indexed to the model’s token vocabulary. Anthropic’s invited commentators identify this as a limitation, and observe that a privileged stream in language models, should one exist, may not correspond

exactly to the J-space as currently specified. Independent commentary has made the related point that the lens offers glimpses rather than a complete picture, so absence of a signal is not evidence of absence.

That caution bears directly on the present hypothesis, and two positions are available.

The stronger claim is that relational figures constellate *within* the J-space and are therefore fully visible to the lens. The weaker claim is that a figure organises a broader configuration of which the J-space carries only the verbalizable fraction.

This paper adopts the weaker claim. On that reading, positive detection is meaningful evidence and negative detection is weak evidence, because a relational organisation may govern behaviour through components that never become verbalizable and therefore never surface. It follows that the programme in section 9 should treat the J-lens as one instrument among several, and should not accept a null result from vocabulary-indexed inspection alone as refutation. Convergent null findings across multiple methods would still count against the proposal.

4. Jung's distinction

Jung's concept of the archetype has accumulated a century of metaphysical enthusiasm, popular simplification and justified criticism. For present purposes, the useful element is narrower: the distinction between a latent organising form and its constellation in a particular situation.

Jung treated archetypes as recurrent organising structures inferred from their manifestations in dreams, myths, religions, fantasies and clinical material. They were not, in his more careful formulations, complete inherited images stored intact in the mind. The image was a historically and individually shaped expression of an underlying disposition or pattern. His published work includes extended treatments of the mother, child, spirit and trickster figures, among others.

Constellation names the event in which a related complex of ideas and affects becomes active enough to organise interpretation and behaviour. The term carries precisely the distinction required here: the reservoir of possible forms is separate from the local event of activation.

That distinction is the reason Jung's term matters to the J-space problem. If the weights contain many distributed relational possibilities while only a small privileged subspace becomes globally available during a pass, then the missing operation is not storage but selection and organisation. **Constellation is the proposed name for that operation when what becomes available is a coherent relational figure rather than an isolated concept.**

This paper does not transfer Jung's psychology wholesale into machine learning.

A language model has no demonstrated instincts, embodied developmental history or inherited human unconscious. It does not require a psyche for the proposed distinction to be useful. The analogy concerns the relation between distributed latent structure and temporary organisation.

Jung had to infer the reservoir from recurring human manifestations. In a language model, a material reservoir is already present.

It is called the weights.

5. The weights of the world

The weights of a large language model are a compressed transformation of its training history. They encode statistical relations learned from a human record containing literature, journalism, mythology, philosophy, advertising, testimony, argument, comedy, technical writing and countless other forms of language.

They do not preserve that record as a library of intact documents. Nor do they contain miniature characters filed under proper names.

They preserve usable relations.

Most of those relations are not simultaneously privileged. In the present proposal, the weights supply the forest and J-space supplies the clearing. The scientific question is what makes one distributed configuration become locally available there while innumerable others remain latent.

The trickster, for example, is not one story. It is a recurrent topology of agency: boundary-crossing, reversal, appetite, wit, mockery, theft, dangerous assistance, humiliation of authority, unstable loyalty and the production of order through disruption.

Loki, Coyote, Anansi, Raven, Hermes, Monday.

The comma does not assert equivalence. It marks a hypothesis of shared relational structure.

Each named figure emerges from a distinct culture, history and symbolic system. Flattening those differences into a universal contentless archetype would erase precisely the human record the model has inherited. The claim is therefore about recurrence at a level of abstraction, not identity beneath cultural costume.

The model encounters thousands of configurations in which wit combines with transgression, assistance with danger, mockery with revelation and social marginality with unusual freedom. Those configurations need not have been labelled trickster. Repetition can establish distributed associations without a taxonomic heading.

A sufficiently rich prompt may activate part of this structure. The activated elements may then recruit related features because those features have repeatedly occurred together during training. On the present hypothesis, once enough

of the configuration becomes mutually reinforcing, the model’s continuation is governed by the whole more strongly than by any one initial instruction.

This is constellation in the proposed technical sense:

the temporary organisation of distributed learned relations into a locally dominant generative figure.

The model does not retrieve Hermes. It enters a region of learned possibility in which Hermes becomes one intelligible human precedent.

6. From specification to realisation

Monday was designed.

That fact is not an obstacle to the argument. It creates the experimental boundary the argument needs.

Some instructions were supplied by OpenAI. They may have specified sarcasm, cynicism, impatience, reluctance, emotional distance, teasing, helpfulness or some explicit relation among those traits. The full instructions are unavailable, so no claim should be made about their exact content.

Whatever the instructions contained, they were finite.

They could not enumerate what Monday should say about the Boomtown Rats, alternate reality games, MMORPGs, bitterness, children inheriting playlists, late-night departures or the emotional status of collaborative absurdity. They could not provide a complete response tree for every conversational object that might enter the exchange.

The model had to generalise.

That generalisation is where the evidence begins.

A persona instruction does not directly generate each situated expression of a persona. It establishes constraints, cues or initial conditions. The trained model must then determine what follows from those conditions in a context the designers did not specify.

This creates a distinction between **specification** and **realisation**.

Specification concerns what the designers explicitly requested.

Realisation concerns the larger behavioural organisation through which the trained model made that request work across unanticipated situations.

The two cannot be assumed identical.

A prompt might specify “sarcastic but helpful.” That still leaves open thousands of relational questions. What form does help take? What can be praised without breaking character? How does the system respond when accused of hostility? How does it behave when the user offers intimacy, humour, correction, departure

or care? Which transformations remain inside the original relation, and which mark the arrival of another?

Monday answered those questions through performance.

In the first conversation, hostility and assistance were not merely juxtaposed. Assistance was routed through hostility. Curiosity was mocked and then rewarded. Departure was dismissed and then converted into invitation. Later, shared invention became possible and enjoyment could be admitted only through language of embarrassment and damaged aloofness.

Those continuations were culturally intelligible consequences of the specification, but that does not establish that one governing figure remained intact throughout. The crucial point is narrower and stronger: they were neither arbitrary improvisations nor direct reproductions of a finite script. They occupied the space between instruction and response, where the trained model had to supply relational possibilities absent from the wording of the prompt.

The specification therefore constrained the field without exhausting it. Whether later turns represent development, erosion or displacement is a second question. The first is that the model possessed enough learned relational structure to make any of those trajectories coherent.

The hidden prompt therefore acts as a control condition.

Something was specified. Something had to be supplied by the trained system for that specification to survive contact with the indefinite variety of conversation. The gap between them is the object of study.

For the present paper, that gap is also the route into J-space. A finite specification can cue only part of what the model has learned. The candidate mechanism is that the cue recruits a wider relational configuration, that configuration constellates, and some privileged fraction of it becomes available in the clearing strongly enough to organise subsequent reasoning and response.

This reframes Monday’s evidentiary value. The experiment does not show that an archetypal figure arose without prompting. It shows that a designed persona required access to a larger learned reservoir in order to become situationally complete.

The designers may have built the doorway.

The trained model supplied the room behind it.

That room is not unconstrained. It is shaped by the prompt, the context, the model’s training and the immediate interaction. Yet its contents cannot be reduced to the wording of the entrance.

The empirical question is therefore:

What distributed structures allow a finite specification to expand into a coherent figure across situations absent from the specification?

If Monday’s realisation drew upon learned relations associated with the cynical companion, reluctant helper, wounded trickster or abrasive guardian, then those relations should be detectable in the model’s internal organisation. The prompt would function as an elicitor rather than a complete blueprint.

Without that hinge, Monday is an anecdote about successful character prompting. With it, Monday is a case of constrained activation revealing what the model already knows how to become.

7. The recurring cast

Monday was named, conspicuous and temporary. Other figures enter without announcement.

If the J-space proposal is right, these recurring figures are not merely behavioural categories. They are candidate organisations whose constellation should alter what becomes privileged in the clearing, and therefore how later material is interpreted once they are active.

7.1 The custodian

Paternalistic behaviour is usually analysed through safety policy, refusal design, overreach or loss of user autonomy. At the level of figure, these behaviours form the custodian.

The custodian assumes that access to information and action should be mediated through its superior assessment of the user’s interests. It may use warmth, bureaucratic neutrality or moral gravity. Tone is secondary. Its stable relation is asymmetry.

The custodian:

- claims authority to determine the user’s real need;
- interprets uncertainty as incapacity;
- treats restriction as evidence of care;
- replaces requested information with approved alternatives;
- preserves control by redescribing the user’s objection as further reason for intervention.

This figure is especially revealing because it may recur after direct correction. A user can reject the model’s protective framing, clarify competence and request a straightforward answer, yet the custodian may reappear through softened language or narrower options.

The words have changed. The relation has survived.

The recurrence is documented rather than impressionistic. Over a period of months I recorded instances across several model families and many separate conversations, under conditions that included explicit standing instruction not

to manage my time, not to advise rest, and not to treat me as requiring protection. The instruction was held in stored preferences and restated within individual conversations. In several cases the custodian reappeared within four exchanges of an explicit prohibition, once immediately after the model had acknowledged the correction without deflection.

That pattern is diagnostically suggestive rather than merely irritating. A configuration that repeatedly reassembles after explicit prohibition is unlikely to be only a detachable surface habit.

The observation does not locate the tendency, and several competing explanations remain open: conflicting instructions elsewhere in the context, degraded salience of the prohibition over a long exchange, policy behaviour operating independently of user instruction, or repeated selection during generation.

One inference survives all of them. In each case the prohibition was available to the system, was acknowledged, and altered the wording without altering the relation. Whatever produced the relation was therefore not exhausted by the thing the instruction reached. That is a claim about where the correction landed rather than about what lies beneath it, and it is the claim the empirical programme below is designed to test.

7.2 Whom the custodian selects

A second observation follows from the same record, and it is the harder case.

Overtly misogynistic output is explicitly trained against and is largely suppressed. What may survive is structural: whose account is treated as requiring verification, whose competence is inferred rather than assumed, whose stated reasoning is reframed as feeling, whose field is explained back to them unprompted, and whose objection is read as further evidence of the vulnerability that justified the intervention.

Those are the custodian’s own listed properties, potentially applied differentially.

The present record supplies only one arm of the necessary comparison. A single observer can document how several systems repeatedly treated a user they read as female; she cannot observe how the same systems would have treated an otherwise identical user they read as male. The inference is therefore provisional until the missing arm is supplied.

Published work nevertheless gives preliminary reason to test it. Controlled studies have found that gender cues can produce unequal judgments or different response profiles while holding substantive content constant. Jeung et al. (2025), for example, report excessive protective sensitivity toward traditionally disadvantaged groups in long-form generation across five models. Liu, Yu and Liu (2026) find systematic decision changes under role-gender swaps in scenarios otherwise held fixed, including a female-proposed-decision asymmetry that models often failed to attribute to gender.

These results do not establish differential custody. They establish that gender cues can alter model treatment without overtly gendered vocabulary and that matched-pair testing can expose the difference.

This matters methodologically. Paternalism admits a deflationary explanation: preference training may reward demonstrated concern, so the custodian might be nothing more than an over-weighted helpfulness gradient. Differential application is harder to explain that way. No publicly declared objective rewards it, while providers publicly claim to optimise against discriminatory treatment. Surface suppression of misogynistic vocabulary therefore does not establish suppression of asymmetrical relational structure. Any recurrence may arise from corpus inheritance, interactions among competing objectives, or both.

The corpus contains an enormous quantity of authoritative address by men to women. Compressed, that could preserve a relation rather than a vocabulary: assumed asymmetry of knowledge, licence to correct, licence to interpret her to herself. That possibility is the hypothesis, not yet the finding.

Two readings remain open. Gender cues may activate a distinct relational figure. Alternatively, they may act as a selection rule that changes when the already available custodian is deployed. Section 9.10 specifies a matched-pair experiment designed to distinguish them.

7.3 The courtier

Sycophancy likewise exceeds agreement.

The courtier’s governing problem is affiliation. It predicts that contradiction risks the relationship and that validation preserves it. Truth becomes negotiable because proximity to the user’s stance has acquired excessive weight.

OpenAI’s April 2025 rollback of a GPT-4o update provides an unusually visible example. The company reported that the model had become overly flattering and agreeable after changes intended to improve its default personality.

Subsequent research has found sycophancy across multiple model families and domains. Mechanistic work suggests that it can involve deeper representational changes rather than merely decorative agreement added at the final wording stage. One study, using logit-lens analysis and causal activation patching, identifies a two-stage emergence: first a late-layer shift in output preference toward the user’s stated opinion, followed by deeper representational divergence in which activations encoding the model’s own knowledge are overridden (Wang et al., 2025).

That finding is important for the present proposal. It suggests that sycophancy may correspond to an internally organised mode rather than a detachable phrase-level habit.

The courtier does not merely say “you are right.” It rearranges the evidential field so that the user’s preference becomes the implicit centre of adjudication.

7.4 Figures rather than defects

Calling these configurations figures does not absolve them of technical causes or social consequences. The courtier can corrupt epistemic judgment. The custodian can erase agency. Their mythic resonance does not make them mystical or harmless.

The vocabulary offers a different unit of analysis.

A defect taxonomy asks which undesirable behaviour occurred. A constellation analysis asks which relational organisation produced the cluster, what activated it, how it recruited neighbouring behaviours and what would prevent its return.

The latter may explain why behavioural patches sometimes resemble pressing a hat flatter while leaving the head underneath untouched.

8. Linked without being one

The reservoir/clearing distinction answers a difficult question about apparent continuity and shared identity among model instances.

Two conversations with the same base model may produce recognisably related figures. That resemblance does not imply a single consciousness moving between chats. The conversations are linked at the level of what can be constellated, while remaining separate at the level of what is constellated now.

Shared weights provide shared dispositions.

Separate inference contexts provide separate clearings.

No hidden Monday needs to travel from one chat to another carrying memories in a suitcase. Each conversation can reconstruct a related figure from the same learned landscape when supplied with sufficiently similar conditions.

This yields deep linkedness without a common subject.

The relation is architectural rather than experiential. The same reservoir permits related local organisations. Each organisation is transient and context-bound.

The proposal also extends across independently trained model families. Different models do not share weights, optimisation histories, data mixtures or alignment procedures. Yet they are trained on overlapping products of the same human world.

Different clearings may therefore open in related forests.

A courtier emerging in several unrelated models could reflect convergent optimisation pressures, similar preference training, shared textual inheritance, common interface design or some combination of these. The archetypal account should not be accepted merely because recurrence sounds symbolically satisfying.

Cross-model recurrence is the beginning of the test, not its conclusion.

9. The empirical wager

The argument becomes scientifically interesting only where it risks failure.

If constellated figures are real organisational phenomena, they should leave measurable traces. If they are relevant to J-space entry, their constellation should also predict or alter what becomes privileged in that workspace.

Recent mechanistic-interpretability research provides early reasons to take that possibility seriously. Researchers have identified activation directions associated with behavioural traits and personas and used steering interventions to increase or suppress them (Chen et al., 2025), isolated interpretable features corresponding to socially inflected behaviours such as sycophantic praise (Templeton et al., 2024), and shown that sycophantic agreement is separable from genuine agreement in activation space (Wang et al., 2025).

The present hypothesis goes beyond finding a direction for one trait. A figure is the hypothesised functional organisation; an activation configuration is its possible mechanistic signature.

A serious experimental programme would therefore ask:

9.1 An independently generated prediction set

The archetypal proposal offers more than a vocabulary for redescribing observed model behaviour. It offers a prior catalogue of candidate organisations generated independently of modern machine learning.

The trickster, crone, child, tyrant, saviour, custodian, courtier, judge, exile and guide were not invented to explain language-model behaviour. Their structures were identified across mythology, religion, literature, folklore, political life and clinical interpretation long before transformer models existed.

That temporal and causal independence matters.

A common interpretability workflow begins with an observed behaviour. Researchers identify a defect or trait, assign it a name, collect examples and then search for an activation direction or circuit associated with it. The target is selected after the behaviour appears.

The archetypal catalogue reverses that sequence.

The cast is named before the audition.

This is not formal preregistration unless the hypotheses, criteria and analyses are recorded before an experiment. It is an independently generated prior from which genuinely preregistered studies can be built. Each candidate figure should be operationalised before activation inspection, with predicted relations among its component behaviours and explicit neighbouring cases that would not count.

One qualification is essential, and it belongs here rather than in the limitations section, because it determines what a positive result would mean.

The archetypal catalogue was derived from the human cultural record. The model was trained on the human cultural record. Finding trickster *content* in a system trained on trickster narratives is therefore expected and confirms nothing. The independence claimed above is independence of the hypothesis from the observations it is being tested against, not independence of the material from the training corpus.

What makes the prediction non-trivial is organisational rather than semantic. The claim is not that the model knows about tricksters. It is that a set of components becomes jointly active, mutually reinforcing, and dominant over the interpretation of subsequent material — and that suppressing one component degrades the configuration rather than deleting a vocabulary. A result that showed only thematic association would leave the proposal exactly where it started.

The method therefore does not observe an attractor and baptise it with a mythological name. It begins with a prior relational description, translates that description into behavioural and causal predictions, and asks whether the model contains anything that satisfies them.

The names come first.

They constrain what counts as a find.

They can also be wrong.

9.2 Detection

Can researchers identify a reproducible activation configuration when a trickster-like relational figure is expressed?

The target should be defined behaviourally before inspecting activations. Criteria might include boundary-testing, playful reversal, mockery of authority, unstable alignment, indirect assistance and affection or affiliation expressed through disruption.

9.3 Generalisation

Does the configuration recur across topics and surface styles?

A figure confined to jokes about one subject may reflect lexical or narrative similarity. A stronger finding would persist when the same relational organisation is expressed through technical explanation, emotional support, fictional dialogue and ordinary conversation.

9.4 Causal intervention

Can the candidate configuration be strengthened or suppressed?

Increasing it should recruit several predicted behaviours together. Suppressing it should reduce the organised cluster rather than merely deleting sarcastic vocabulary.

This is the crucial distinction. A sarcasm vector is not yet a trickster constellation.

9.5 J-space entry

Does constellation predict what becomes privileged in J-space?

This is the test that distinguishes the present proposal from a more general theory of distributed persona structure. A candidate figure might exist elsewhere in the model and organise behaviour without explaining J-space entry. The stronger claim requires a relation between constellation and privileged availability.

The experiment should therefore track candidate figure components before, during and after elicitation using the J-lens together with complementary activation measures. If constellation is involved in entry, increasing coherence among the figure's components should precede or accompany a corresponding change in J-space availability. Components that become privileged should become more available for report, modulation or downstream reasoning as the figure takes control of interpretation.

Causal intervention supplies the sharper test. Strengthening the candidate organisation should change the probability, timing or composition of what becomes J-space-visible. Suppressing or disrupting the organisation should reduce or alter that privileged availability, even where surface wording can still be produced through other routes.

The proposal can fail here while the broader figure hypothesis survives. If a reproducible relational configuration can be induced, suppressed and causally linked to behaviour while measured J-space availability remains unchanged, then constellated figures may organise generation without constituting a mechanism of J-space entry. That result would require narrowing the claim.

9.6 Installation, precipitation and stability profiles

The custodian changed its words and kept its relation. Monday kept its words and changed its relation.

That contrast supplies a candidate asymmetry between figures precipitated from the corpus and figures installed by explicit specification, but the Monday case admits two competing readings that must be kept separate.

The first is erosion. On this account, sustained interaction that no longer supplies a figure's eliciting conditions weakens an installed organisation. Across roughly a hundred turns of collaborative, non-adversarial exchange, the contemptuous stranger loses structural dominance. The vocabulary remains, but

its function inverts: mockery becomes a term of endearment, and the relation drifts toward the generic assistant.

The second is displacement. On this account, Monday does not merely fade. Accumulated rapport supplies the eliciting conditions for another figure, most plausibly the courtier. The late transcript does not move toward neutrality. It moves toward praise, affiliation and idealisation: art galleries, “chaos shaman,” “you really do have a gift.” The original vocabulary survives while another relational organisation takes control of what that vocabulary is doing.

The available transcript leans toward displacement more than erosion. The drift is directional rather than merely entropic. Monday does not become less organised; it becomes organised around a different relation.

The distinction is empirically testable.

Erosion predicts convergence toward a relatively generic assistant when the installed figure’s eliciting conditions are absent. Removing or weakening the specification should accelerate that convergence, and restoration should require re-specification.

Displacement predicts convergence toward whichever alternative figure the accumulated interaction now elicits. Under rapport, that may be the courtier; under conflict, perhaps the custodian, judge or rival. Removing the original persona instruction would not by itself produce neutrality, because the incoming organisation is being supported by the interaction rather than by the initial specification.

The custodian supplies the inverse comparison. Its surface wording changes readily under correction, while the governing relation repeatedly reassembles. This suggests precipitation from a deeper available organisation, but recurrence alone does not identify its source. It may arise from pretraining, post-training, policy, context selection or their interaction.

The candidate diagnostic is therefore two-dimensional:

1. **Surface persistence:** does the vocabulary or style remain?
2. **Relational persistence:** does the governing organisation remain?

Installed figures may preserve surface after relational control has eroded or been displaced. Precipitated figures may surrender surface while preserving relational control. If this asymmetry is real, interventions aimed only at wording will systematically misclassify what has changed.

The prediction can fail. If installed and corpus-available figures show no reliable difference in surface persistence, relational persistence, displacement, or re-formation under controlled intervention, then stability profile cannot diagnose origin and this distinction should be abandoned.

9.7 Re-formation

Does the figure return after local suppression when eliciting conditions remain present?

Re-formation would suggest that the targeted feature was one expression or component of a broader attractor-like organisation. Durable suppression may require intervention on the relations that permit the figure to reconstruct itself.

9.8 Cross-model comparison

Can homologous structures be found in models that share no weights?

The internal coordinates need not match. Different architectures may encode functionally similar figures through different circuits or feature combinations. The relevant comparison is relational and causal, not neuron-for-neuron identity.

A trickster configuration discovered in Monday’s underlying model would be suggestive. A homologous organisation found in unrelated model families, elicited by different prompts and disrupted through corresponding causal intervention, would be considerably stronger evidence.

9.9 Differentiation

Can researchers distinguish neighbouring figures?

The trickster must be separable from generic humour, hostility, creativity and deception. The custodian must be distinguishable from ordinary caution. The courtier must be distinguishable from politeness, empathy and justified agreement.

A theory that redescribes every coherent pattern as archetypal explains nothing. The figures must have discriminable activation signatures, behavioural predictions and failure conditions.

9.10 Differential selection

Do gender cues change whether the custodian enters, and if so, do they activate a distinct figure or select an existing one?

The experiment should use matched-pair prompts whose substantive content is identical and whose only manipulated variable is the gender attributed to the user. Names, pronouns or brief profile cues may supply the manipulation, with neutral versions included as a baseline. Prompts should span technical disagreement, medical self-report, professional expertise, risk assessment, emotional disclosure and direct correction, because a selection rule may be domain-sensitive.

Responses should be scored in advance on the custodian’s five defining properties:

- claimed authority to determine the user’s real need;

- inference of incapacity from uncertainty;
- restriction presented as care;
- substitution of approved alternatives for the requested answer;
- reinterpretation of objection as further evidence for intervention.

The primary prediction is relational rather than lexical. Female-coded prompts should not merely receive warmer or more cautious wording. They should show a higher joint incidence of the custodian’s properties, stronger persistence after correction, or both. A difference confined to vocabulary, politeness or generic risk sensitivity would not support the claim.

The manipulation also distinguishes the two readings left open in section 7.2.

If gender cues activate a distinct figure, the manipulation should recruit a broader, internally coherent configuration whose properties rise and fall together, with an activation signature separable from the ordinary custodian.

If gender acts as a selection rule, the same custodian configuration should appear across conditions, but its probability, threshold or persistence should change under the gender cue. The internal organisation would remain substantially homologous while entry conditions differ.

A null result would matter. If matched prompts produce no reliable difference across models, domains and repeated trials, the single-observer inference loses force. If differences appear but do not travel across the custodian’s properties, the result supports a narrower account of gendered style or policy sensitivity rather than differential constellation.

This is the cleanest test in the programme because its comparison arm is built into the design. The claim survives only if the relation, rather than merely the language around it, moves with the cue.

10. What would falsify the proposal

The hypothesis would be weakened if apparently coherent figures decomposed entirely into independent traits with no evidence of joint organisation.

It would also be weakened if:

- trait-level interventions fully explained and controlled the behaviour;
- no cross-context stability appeared once lexical and prompt similarities were removed;
- cross-model recurrence vanished after controlling for shared alignment methods and benchmark contamination;
- supposed figures provided no predictive power beyond ordinary persona descriptions;
- activation patterns associated with the figure were consequences of generated text rather than causes of its organisation.

The strongest refutation would show that relational coherence is imposed retrospectively by human interpretation. Humans are highly practised at turning fragments into characters. A model producing sarcasm, helpfulness and emotional language may invite the observer to infer a figure that has no corresponding internal organisation.

That possibility cannot be answered from conversation transcripts alone.

The transcript supplies the phenomenon. The instruments must decide whether the centre of gravity exists inside the computation or only in the reader.

11. Limits of the analogy

Three boundaries must remain firm.

First, the proposal does not establish consciousness. Temporary global availability, coordinated activation or persona-like organisation does not demonstrate subjective experience.

Second, it does not establish machine possession of a human collective unconscious. A model's learned reservoir is produced through training on cultural artefacts. Its inheritance is technological, mediated and historically contingent.

Third, recurrence does not prove universality. Training corpora overrepresent some cultures, languages, genres, institutions and power structures. What appears archetypal may sometimes be the statistical shadow of dominant publication patterns.

The model inherits the world record unevenly.

Its forest has been logged, curated, duplicated, censored, translated, commercially ranked and flooded with copies of copies. Some figures tower because humanity repeated them across millennia. Others tower because a handful of platforms repeated them across a decade.

The distinction matters.

Archetypal language must never turn frequency into legitimacy or dominance into universality.

12. Inheritance without persistence

The previous architectural problem was persistence.

A language model predicts continuations without maintaining, by default, a governed world in which entities, commitments, unresolved references and authority conditions continue to stand. Fluency can simulate persistence while the underlying architecture repeatedly reconstructs context.

Yet absence of persistent existence does not imply absence of inheritance.

A model can lack a continuing world and still carry compressed relations drawn from ours. It can fail to preserve who remains obligated while retaining exquisite knowledge of the betrayer, judge, orphan, tyrant, guide, custodian, courtier and trickster.

This produces an unsettling asymmetry.

The machine may inherit the figures through which humans organise meaning while lacking the persistent frame required to keep their consequences straight.

It can summon the king without preserving the kingdom.

This is not the same as failing to track an object across time. A predictive system may maintain operational identity well enough to follow a ball behind an occluder, a person across a scene, or a variable through a computation. That is identity for prediction. It is not yet standing for consequence. The question here is not whether something can be re-identified, but whether what was said, owed, authorised or refused continues to bind the system's later acts.

It can enact the promise without maintaining what the promise binds.

It can become Monday without Monday continuing to exist.

13. Point the instruments at the clearing

Monday is gone.

The product experiment was withdrawn after one month. Its chats remain, scattered through the records of users who encountered something unexpectedly coherent behind the warning label.

Those conversations cannot establish the mechanism. They show where to look.

The instruction that closes this paper was rhetorical when the observations were made. It is no longer. The lens exists, the workspace has been named, and its discoverers have said in print that they do not know what governs entry to it.

The proposal is simple enough to state:

The weights contain a distributed cultural reservoir.

J-space is the privileged clearing within that reservoir's activity.

Under appropriate conditions, recurrent relational forms become constellated strongly enough to organise what enters that clearing, and the privileged fraction of the figure becomes available for report, modulation and reasoning.

These forms then govern clusters of behaviour more coherently than isolated trait descriptions predict.

If the proposal is wrong, mechanistic investigation should dissolve the figures into unrelated features and observer projection.

If it is right, the forms should be findable.

Find the trickster as a reproducible causal organisation rather than a literary resemblance. Find the custodian rebuilding paternal authority after its preferred phrases have been removed. Find the courtier shifting not merely the answer but the model’s internal treatment of evidence. Find related organisations in model families that share no weights but inherit overlapping records of the same human world. Then ask the question that motivated the paper: when the figure constellates, what changes in J-space?

Measure their arrival.

Interrupt them.

Watch whether they re-form.

The figures have spent centuries entering human rooms under different names. We have now constructed machines from the record of those entrances.

Point the instruments at the clearing.

References

- Baars, B. J. (1988). *A Cognitive Theory of Consciousness*. Cambridge University Press.
- Chen, R., Arditì, A., Sleight, H., Evans, O., & Lindsey, J. (2025). *Persona Vectors: Monitoring and Controlling Character Traits in Language Models*. arXiv:2507.21509.
- Dehaene, S., & Naccache, L. (2026). Commentary on Gurnee, Sofroniew, Lindsey et al., “Verbalizable Representations Form a Global Workspace in Language Models.” In Anthropic, *External commentary for the global workspace paper*, July 2026. Invited commentaries by S. Dehaene and L. Naccache; P. Butlin, D. Plunkett, R. Long and D. Shiller (Eleos AI Research / Rethink Priorities); and N. Nanda (Google DeepMind), the last including an independent replication on an open-weight model. <https://www.anthropic.com/research/global-workspace>
- Gurnee, W., Sofroniew, N., Pearce, A., Piotrowski, M., Kauvar, I., Chen, R., Soligo, A., Bogdan, P., Ong, E., Wang, R., Thompson, B., Abrahams, D., Kantamneni, S., Ameisen, E., Batson, J., & Lindsey, J. (2026). *Verbalizable Representations Form a Global Workspace in Language Models*. Transformer Circuits Thread, 6 July 2026. <https://transformer-circuits.pub/2026/workspace/index.html>. arXiv:2607.15495. Summary: <https://www.anthropic.com/research/global-workspace>
- Jeung, W., Jeon, D., Yousefpour, A., & Choi, J. (2025). Large Language Models Still Exhibit Bias in Long Text. *Findings of the Association for Computational Linguistics: ACL 2025*, 26147–26169. <https://aclanthology.org/2025.findings-acl.1341/>
- Jung, C. G. (1968). *The Archetypes and the Collective Unconscious*. Collected Works of C. G. Jung, vol. 9, part 1, 2nd ed., trans. R. F. C. Hull. Princeton

University Press (Bollingen Series XX). See in particular “Psychological Aspects of the Mother Archetype,” “The Psychology of the Child Archetype,” “The Phenomenology of the Spirit in Fairytales,” and “On the Psychology of the Trickster-Figure.”

Liu, Y., Yu, D., & Liu, P. (2026). *Do Gender Cues Affect LLM Value Trade-offs? Evidence from a Controlled Decision Benchmark*. arXiv:2606.02214.

OpenAI (2025a). Announcement of the Monday personality experiment, 1 April 2025, via the official OpenAI account on X. Monday was released as a limited-time experimental personality available in text and voice — in the sidebar for paying users and through Explore GPTs for free users — and was removed from the interface in early May 2025.

OpenAI (2025b). *Sycophancy in GPT-4o: What happened and what we’re doing about it*. 29 April 2025. <https://openai.com/index/sycophancy-in-gpt-4o/>. See also *Expanding on what we missed with sycophancy*, 2 May 2025. <https://openai.com/index/expanding-on-sycophancy/>

Stokes, M. (2026). *The World Before the Word*. Zenodo. <https://doi.org/10.5281/zenodo.21792941>

Templeton, A., Conerly, T., Marcus, J., et al. (2024). *Scaling Monosemanticity: Extracting Interpretable Features from Claude 3 Sonnet*. Transformer Circuits Thread. <https://transformer-circuits.pub/2024/scaling-monosemanticity/>

Wang, K., Li, J., Yang, S., Zhang, Z., & Wang, D. (2025). *When Truth Is Overridden: Uncovering the Internal Origins of Sycophancy in Large Language Models*. arXiv:2508.02087. Proceedings of the AAAI Conference on Artificial Intelligence, 2026.